

SWEN90004

Modelling Complex Software Systems

Wrapping up

Nic Geard

Review Lecture

Semester 1, 2019

Topics: Concurrency

Concurrency is a design principle for structuring programs to reflect **potential** parallelism.

Compared to a sequential system, a concurrent system is characterised by **multiple** threads of control (sometimes referred to as processes).

These processes need to:

- **communicate**: either via shared memory (Java) or message passing (FSP, Go)
- **synchronise**: ie, one thread can't proceed until another is in a certain state

Development and debugging are difficult due to **nondeterminism**.

Topics: Modelling concurrency

When reasoning about concurrent systems, we **abstract** away absolute time and consider relative timing: the order in which **atomic events** occur.

Atomic events may be **arbitrarily interleaved** (occur in any order).

We typically want to ensure that concurrent programs behave correctly for **any** arbitrary interleaving.

Topics: Interference

We illustrated the problems that can arise when concurrent processes **interfere** with each other; eg if two processes each try to execute the following atomic events:

- read shared value x
- update value $x = x+1$
- write shared value x

Topics: Mutual exclusion

We studied a simple **mutual exclusion** scenario in which two threads each alternated between a non-critical and a critical section of code (in which the shared value was accessed).

Solutions needed to avoid:

- **deadlock/livelock**: when one or more threads want to enter their CS, one must succeed
- **starvation**: a thread wanting to enter its CS must eventually succeed

Topics: Concurrency in Java

Java handles concurrency using threads, created by extending `Thread` or implementing `Runnable`.

A Java thread that is alive can be in one of several states:

- runnable (ready to be executed)
- running (currently being executed)
- non-runnable (eg, sleeping, waiting on another thread, etc)

Topics: Concurrency in Java

We indicate to the Java VM that a thread should be executed by calling `start()`. This causes the VM to execute the thread's `run()` method in its own concurrent thread. The thread will stop when `run()` finishes.

Calling `t.join()` suspends the calling thread until `t` has finished running.

Threads can be interrupted while non-runnable, in which case they will return to being runnable and throw an `InterruptedException`.

Topics: Monitors

A monitor is a construct that handles synchronization.

It provides threads with:

- mutual exclusion in access to shared (private) data
- the ability to wait for a certain condition to become true
- (and to be notified when that condition becomes true)

While a thread has (exclusive) access to a monitor, it is said to hold that monitor's lock.

Topics: Synchronisation in Java

Any object in Java can act as a monitor; all objects have:

- an intrinsic **lock**, which is obtained by a thread when it calls a `synchronized` method, and released when that method returns
- a means of handling which threads are currently waiting to obtain that lock

We can mark a method as a critical section by declaring it with the `synchronized` keyword.

Topics: Synchronisation in Java

Methods relevant to synchronisation in Java:

- `wait()`: Causes the current thread to wait until another thread invokes the `notify()` method or the `notifyAll()` method for this object.
- `notify()`: Wake up a single (arbitrary) thread waiting on **this object's** lock.
- `notifyAll()`: Wake up all objects that are waiting on this object's lock.

Topics: Synchronisation in Java

`wait()` should be wrapped within a `while` loop that checks on the condition it is waiting for:

```
class Account {  
    public synchronized void withdraw(int amount) {  
        while (balance < amount) {  
            try {  
                wait();  
            } catch (InterruptedException e) {}  
        }  
        ...  
    }  
    ...  
}
```

Topics: Semaphores

Manages access to a shared resource; consists of:

- some number (v) of currently available permits
- a wait set (W) of processes waiting for permission

`S.wait()`

`S.signal()`

```
if S.v > 0
    S.v--
else
    S.W = S.W U {p}
    p.state = blocked
```

```
if S.W == { }
    S.v++
else
    choose q from S.W
    S.W = S.W \ {q}
    q.state = runnable
```

Topcis: Concurrency in FSP

FSP is an algebraic language for **modelling** concurrent systems.

An FSP process is defined in terms of sequences of actions.

FSP processes are equivalent to LTS diagrams.

Benefits of formal approaches such as FSP include:

- unambiguous communication about system structure and dynamics
- automated checking of safety and liveness properties

Topics: Checking safety and liveness in FSP

Two types of property that are of interest in concurrent systems:

- **Safety properties**: nothing “bad” (eg, deadlock or interference) ever happens during execution
- **Liveness properties**: something “good” eventually happens during execution (eg, all processes trying to access their critical section eventually do so)

If we model a concurrent process using FSP, we can automatically check that a defined property holds; that is, is true for every possible trace/execution of the model.

Topics: Message passing

The **message passing** paradigm for concurrency removes the concept of shared memory (and associated issues).

Rather, concurrent processes communicate with one another by sending and receiving messages.

Message passing can be **synchronous** or **asynchronous**.

Synchronous communication behaves similarly to a shared action in FSP: both the sender and receiver must occur at the same time.

Asynchronous communication removes the need for a receiver to be ready to receive before a sender is able to send.

Topics: Concurrency in Go

Go uses a message passing approach to concurrency (although it does also allow shared memory).

Concurrency entities in Go are:

- **go-routines**: are very lightweight processes/threads
- **channels**: named entities that allow data to be sent and received between go-routines